

Lagrangian Self-Tuning (λ) over Training

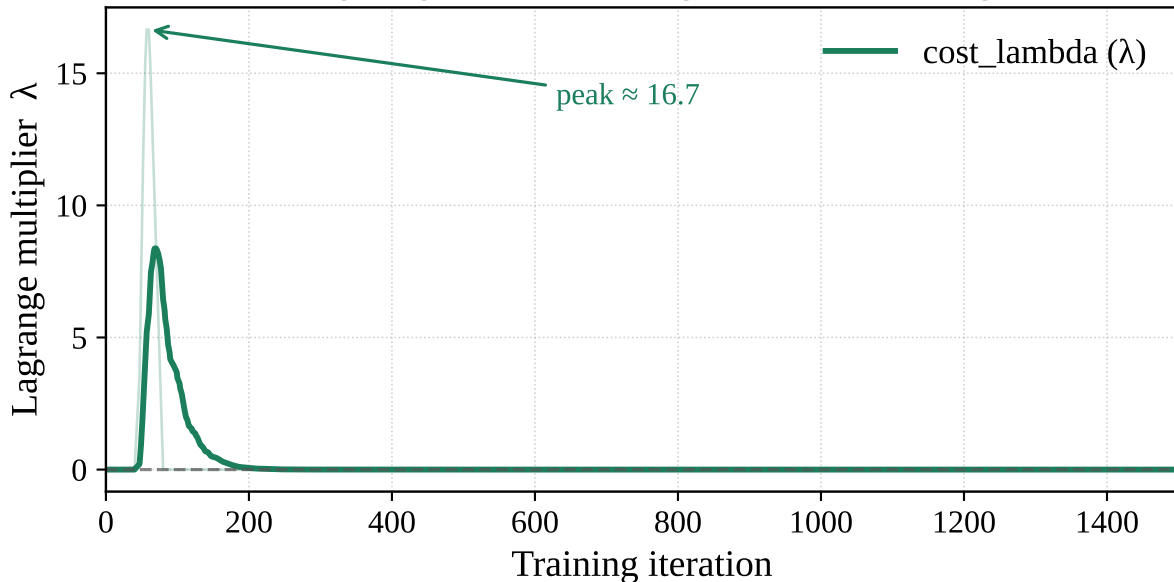


Figure 3. Self-tuning of the dual variable λ (cost_lambda). λ rises to a 16.7 peak to penalise near-singular poses, then relaxes to zero once the constraint is satisfied — never approaching its cap of 100.